

Artificial Intelligence and Decision Systems

Inteligência Artificial e Sistemas de Decisão – MEEC, Maer,

Luís Manuel Marques Custódio

North tower – ISR / IST

lmhc@isr.ist.utl.pt

Problem solving by searching

Often we don't know the algorithm to solve a problem; we only have a specification of what is a solution

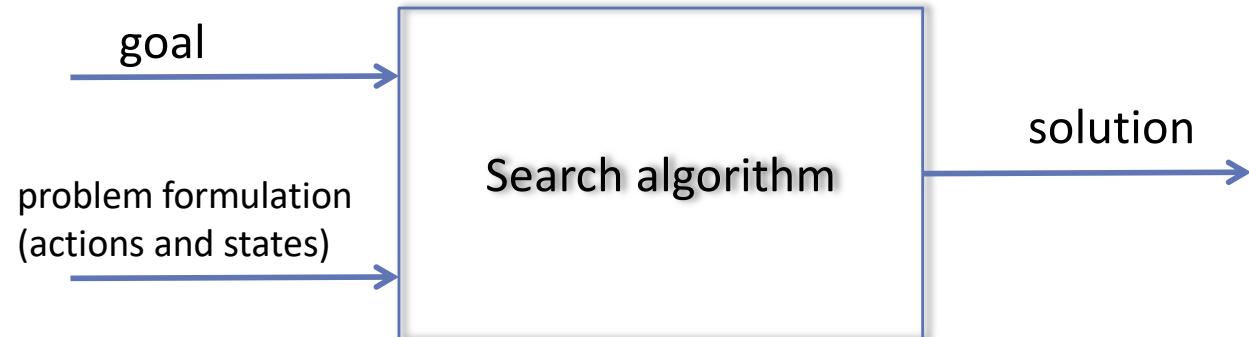
in this case, **we have to search for the solution**

A **problem-solving agent** is a goal-based agent that acts on the environment, leading him to go through a series of states in order to achieve the desired goal

Problem solving is often formulated as a combinatorial **search** problem – agent must find the *right* sequence of actions to achieve the goal

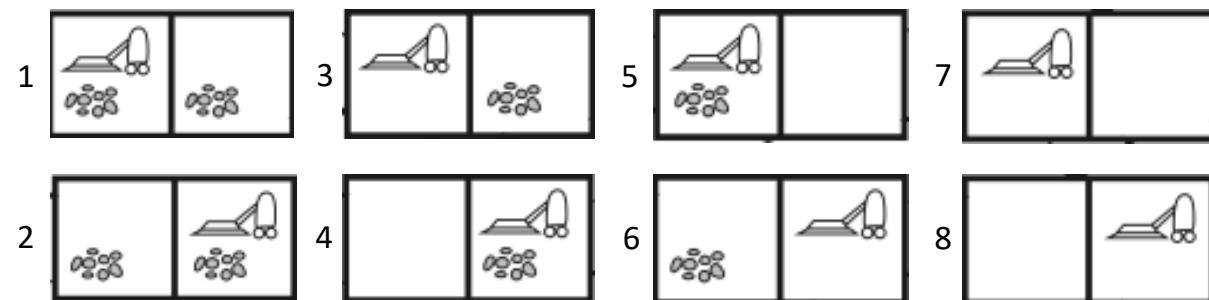
Steps to formulate a search problem

- **Goal formulation**
 - a goal is a (world) state or a set of states
- **Problem formulation**
 - Choosing what actions and states to consider, given the goal
 - an action is a transition between states
- **Search**
 - Process of looking for a sequence of actions that reaches the goal
 - That sequence is called a **solution**
- **Execution**
 - Practical implementation of the solution



Example: vacuum cleaner world

- Goal: cleaning of the two squares
- Actions: {GoLeft, GoRight, Suck}
- States: how many?

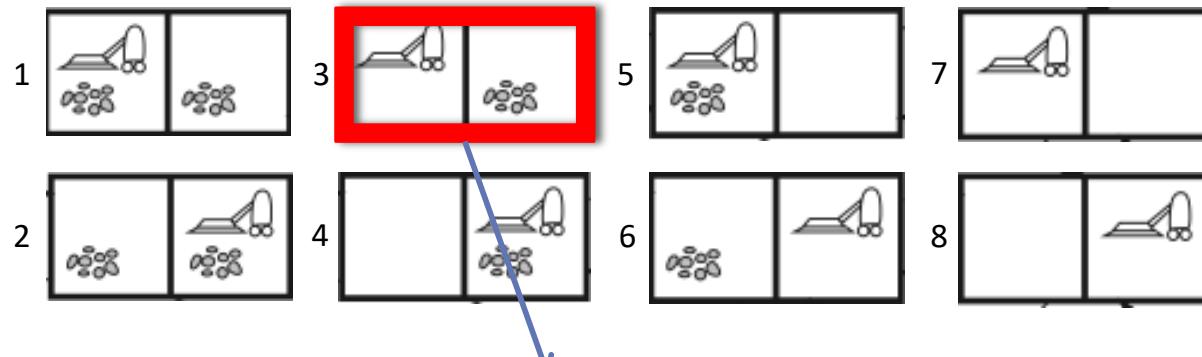


Types of search problems

- Single-state problem

when each action generates a single state transition

world must be fully observable and deterministic



Solution: [right, suck]

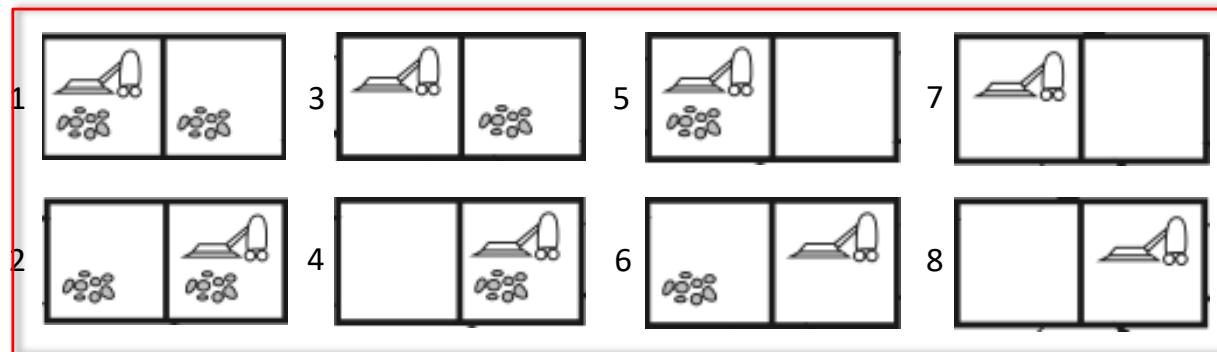
what happens if the world is not fully observable (or non deterministic)?

Examples:

- Lack of sensors
- Faulty cleaner

Types of search problems

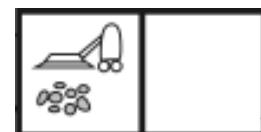
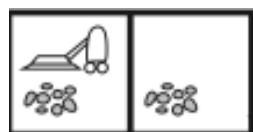
- Multi-state problem
when an action establishes a transition to a set of states



Solution: [right, suck, left, suck]
if the cleaner is ok

What if there are no sensors and the cleaner is faulty?

Example: suppose that the world is in one of these states, there is no dirt sensor, and the vacuum cleaner is faulty: it cleans if there is dirt, but it leaves dirt if it is clean



Which sequence of actions allows
to reach the goal?

Types of search problems

- **Contingency problems**

when the problem involves unpredictable situations whose solution can be selected only when situations occur

Example: {Suck, GoRight, Suck only if there is dirt} [if the robot has local sensing]

- **Exploration problems**

when the agent has no information about the environment or the consequences of their actions

The agent is forced to explore the world in order to build a map of it and learn the effects of their actions

In this course, we only study the single-state problems

Elements of a search problem:

- **Full state space:** set of possible states of the world
- **Initial state:** initial condition of the world
- **Operators:** set of actions that change the world state
 - Path: sequence of states generated by a sequence of operators
 - State space of the problem: set of all states reachable from the initial space by any sequence of actions
- **Goal test:** condition that has to be satisfied at a goal state
 - can be single state, multiple states or defined by a characteristic property of the environment
- **Path cost:** cost associated to a sequence of states/operators
 - it is defined as the sum of the costs of the operators used to form the path

Solution: is an action sequence (a path) that leads from the initial state to a goal state

Elements of a search problem:

- Evaluation of the solution quality:

Path cost

Search cost

in terms of computation time and memory

Optimal solution: is a solution that has the lowest path cost among all solutions

- Problem description

level of detail that the information about the problem must have both in describing states and actions

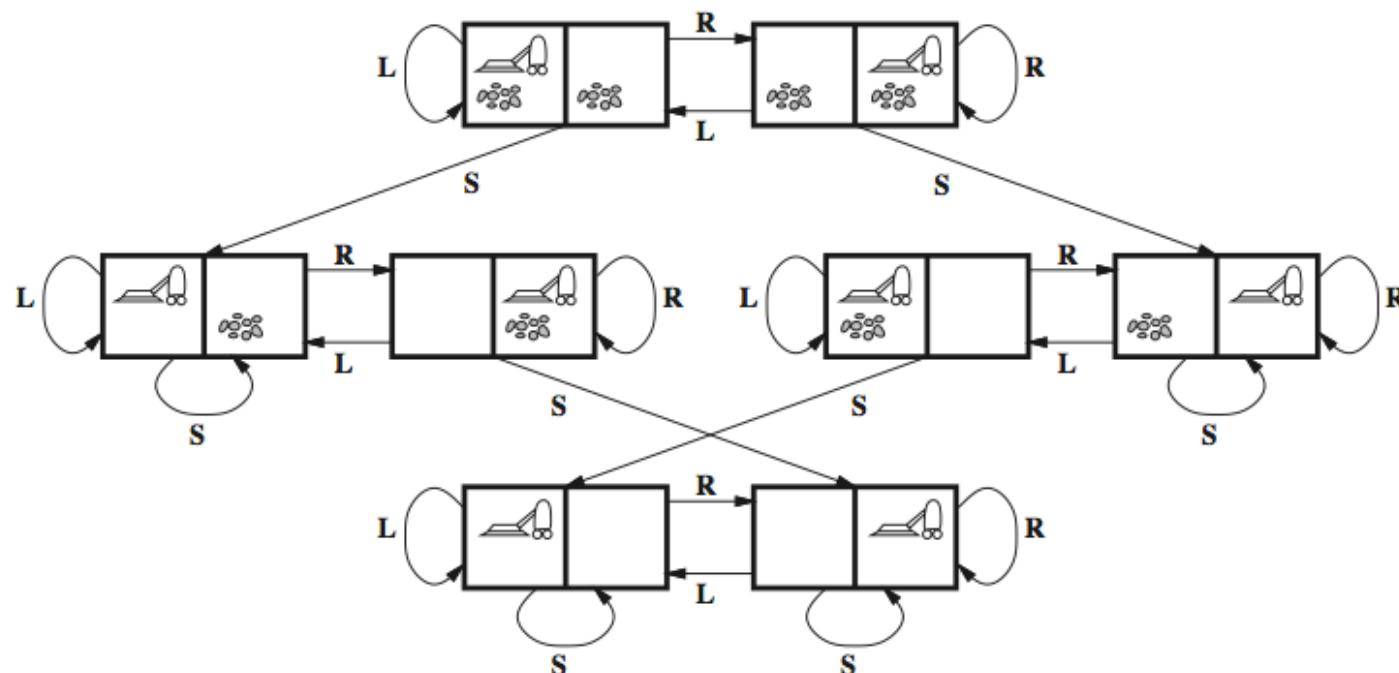
Example: vacuum cleaner

- brand
- color
- wire Length
- energy consumption

Abstraction: process of removing from a problem representation details considered irrelevant for the goal

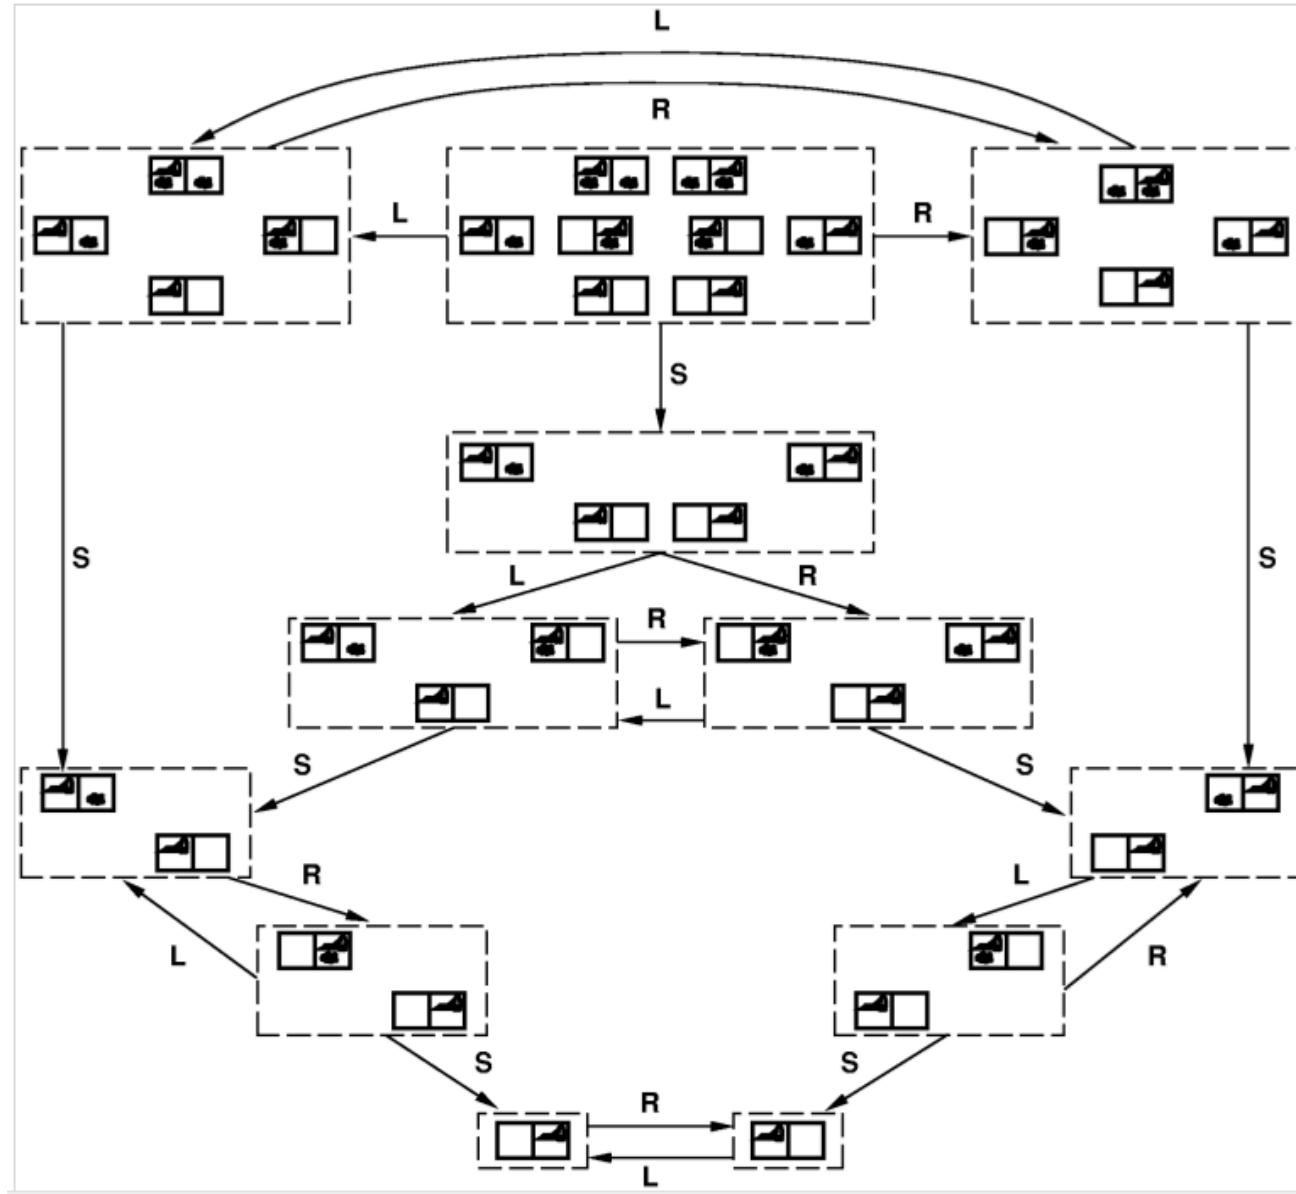
Vacuum cleaner

- states: $\langle r, d_1, d_2 \rangle$, where r is the robot position, d_1 and d_2 are T/F representing presence of dirt in each room
- operators: go left (L), go right (R), and suck (S)
- goal test: $d_1=F$ and $d_2=F$
- initial state: $\langle r, d_1, d_2 \rangle = \langle 1, T, T \rangle$
- path cost: 1 for each operator (also known as step cost)



Search problem examples

Vacuum cleaner with no localization sensors



Search problem examples

8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

States: specify in a 3x3 matrix the location of the 8 numbered tiles plus the blank one

Initial state: any puzzle configuration

Operators: move “blank” to the left, right, up and down

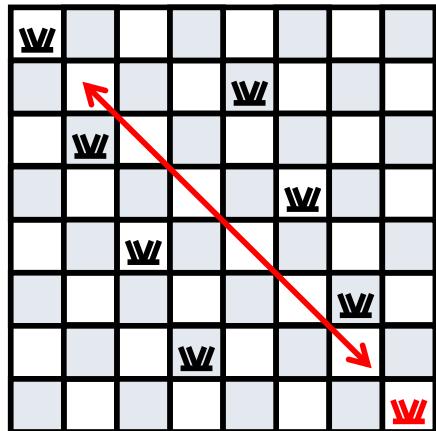
Goal test: checks whether the state matches the goal configuration

Path cost: each step costs 1 (so path cost = number of steps)

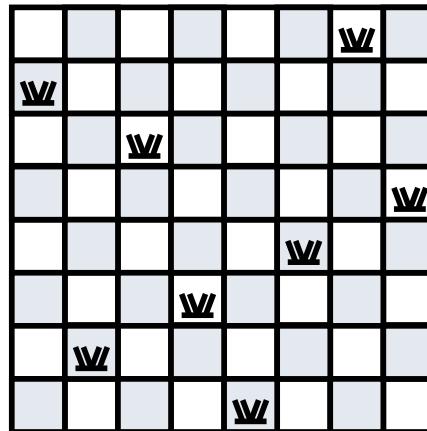
[12]

Search problem examples

8-queens



almost solution



solution

States: location of the eight queens in a chess board

Initial state: empty board

Operators: add a queen to any empty square

Goal test: eight queens are on the board, none attacked

In this case, the number of possible states is $64 \times 63 \times 62 \times 61 \times 60 \times 59 \times 58 \times 57$

$$64^8 = 281\,474\,976\,710\,656$$

8-queens: alternative formulation

States: all possible arrangements of the queens, one per column, with no queen attacking another

Operators: add a queen to any square in the leftmost empty column such that it is not attacked by any other queen

In this case, the number of possible states is 2057 (**check this out!**)

Cryptography

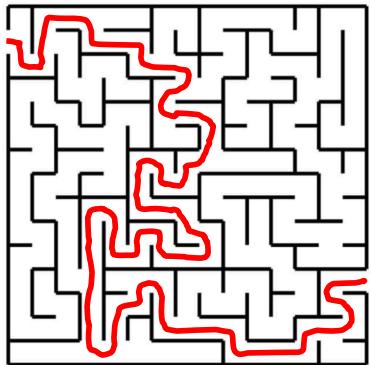
$$\begin{array}{r} \text{T} \quad \text{R} \quad \text{E} \quad \text{S} \\ + \quad \text{D} \quad \text{O} \quad \text{I} \quad \text{S} \\ + \quad \quad \quad \text{U} \quad \text{M} \\ \hline \text{M} \quad \text{E} \quad \text{I} \quad \text{I} \end{array} \qquad \begin{array}{r} 3 \quad 5 \quad 2 \quad 1 \\ + \quad 4 \quad 6 \quad 0 \quad 1 \\ + \quad \quad \quad 7 \quad 8 \\ \hline 8 \quad 2 \quad 0 \quad 0 \end{array}$$

States: a set of letters, some of them associated with a number

Operators: substitute all occurrences of one letter by an unused number

Goal test: all letters were substituted and the numbers satisfy the arithmetic operation

Mazes



States: maze configuration plus current location

Initial state: any maze location

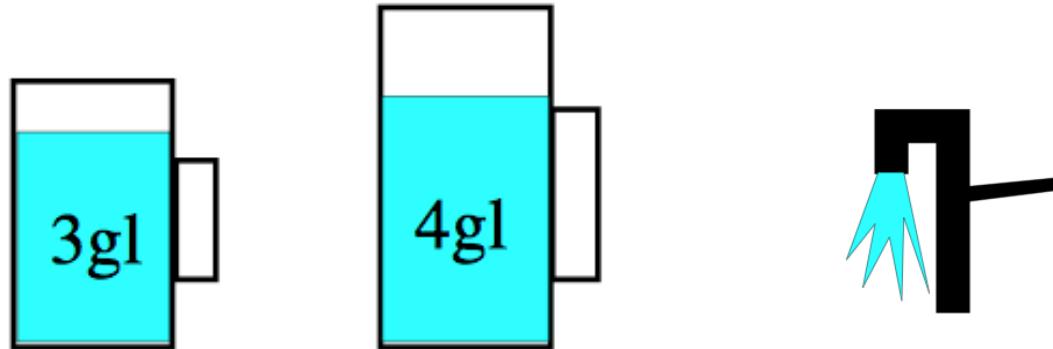
Operators: move to an adjacent and linked position

Goal test: current location = maze exit

Path cost: number of steps

Search problem examples

Water jugs



States: amount of water in each jug (e.g., $<1,4>$)

Initial state: $<0,0>$

Operators: fill J4; fill J3;
empty J4; empty J3;
pour water from J4 into J3 until J3 is full or J4 is empty;
pour water from J3 into J4 until J4 is full or J3 is empty;

Goal test: $<0,2>$

Path cost: amount of water used

Hanoi towers



States: location of the disks in the three poles

Initial state: all disks correctly arrange in the left pole

Operators: move a free disk to another pole, which is empty or all disks in it are bigger

Goal test: all disks correctly arrange in the right pole

Path cost: number of disk moves

Missionaries and cannibals



States: location of the missionaries, cannibals and the boat

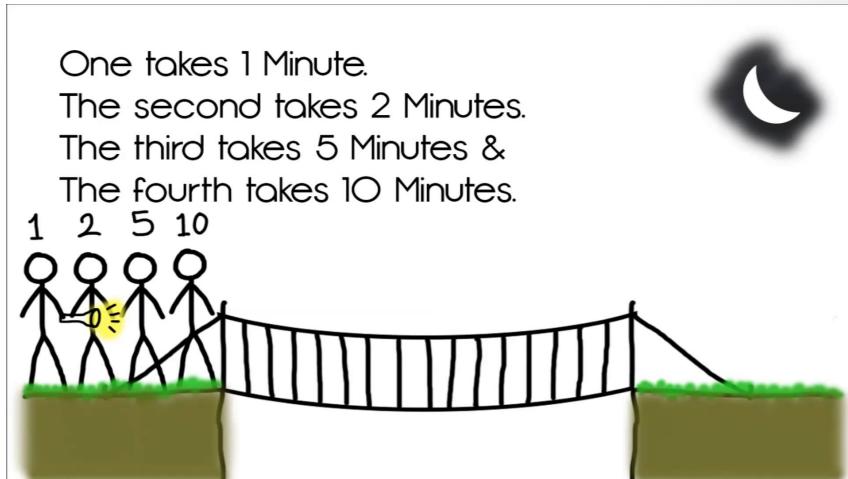
Initial state: all on one side of the river

Operators: move one or two persons to the other side

Goal test: all on the other side of the river

Path cost: number of boat moves

4 Men and the Bridge



States: location of the men and the flashlight

Initial state: all on one side

Operators: move one or two men, with the flashlight, to the other side

Goal test: all on the other side

Path cost: number of minutes needed

Travelling salesman problem

Given a list of cities and the distances between each pair of cities, what is the shortest possible route that visits each city exactly once and returns to the origin city?



Travelling Salesman

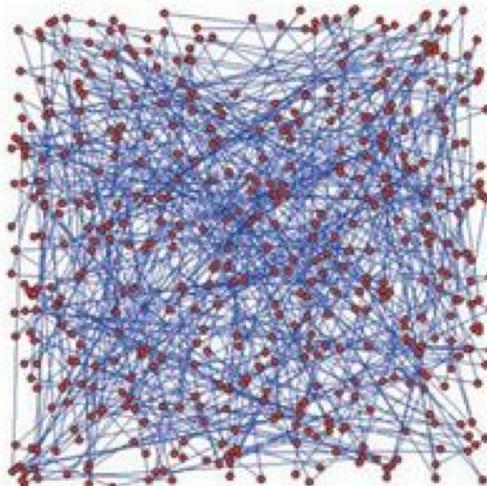
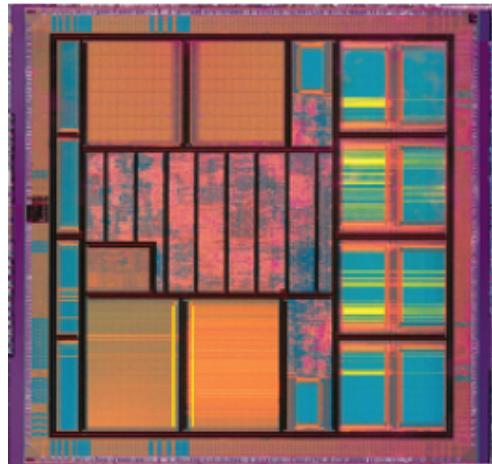


Illustration 3: Path through network with 500 nodes

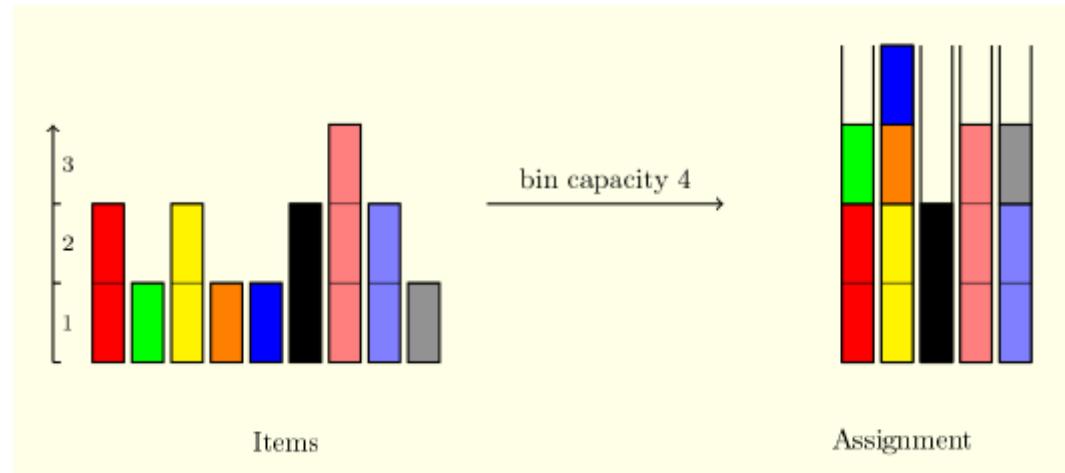
VLSI layout



How to position millions of components and connections on a chip to minimize area, minimize circuit delays, minimize stray capacitances, and maximize manufacturing yield?

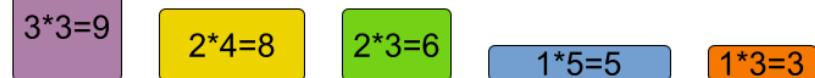
Also, the packing and cutting problems...

Bin packing problem

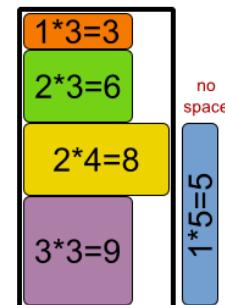


Bin packing

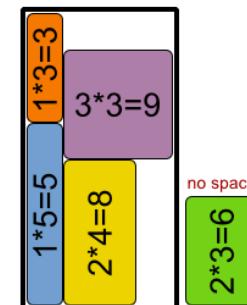
Place each item on a location in a container.



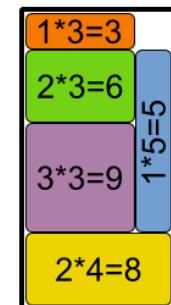
Largest size
first



Largest side
first

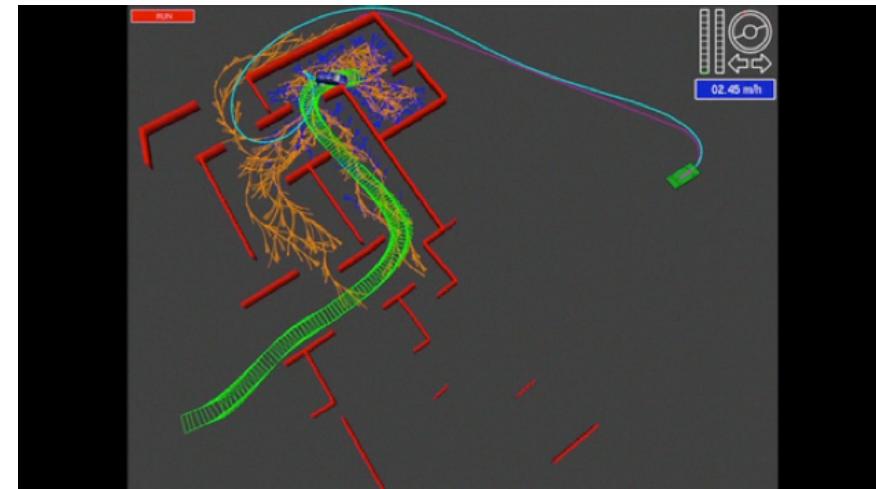
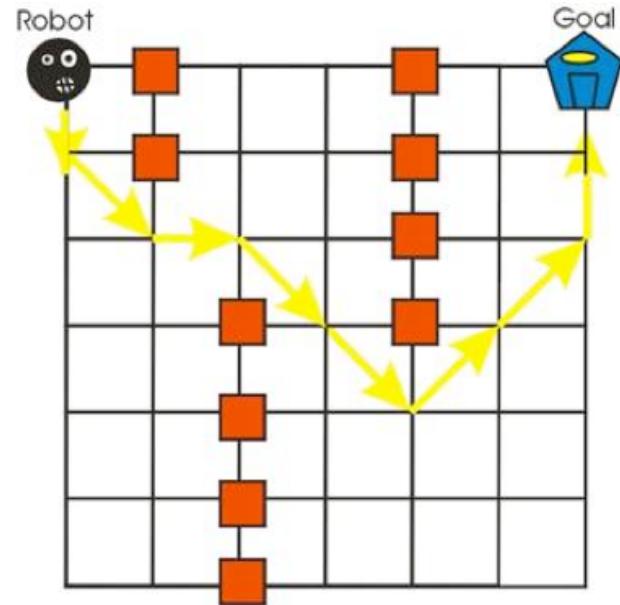


Drools
Planner

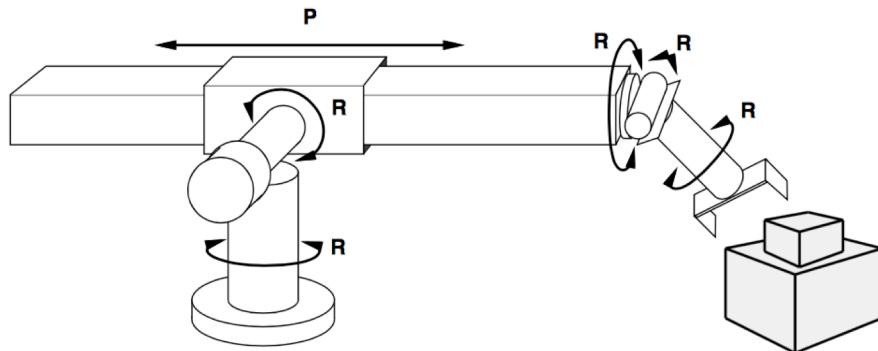


Search problem examples

Robot navigation



Robotic assembly



States: real-valued coordinates of robot joint angles and parts of the object to be assembled

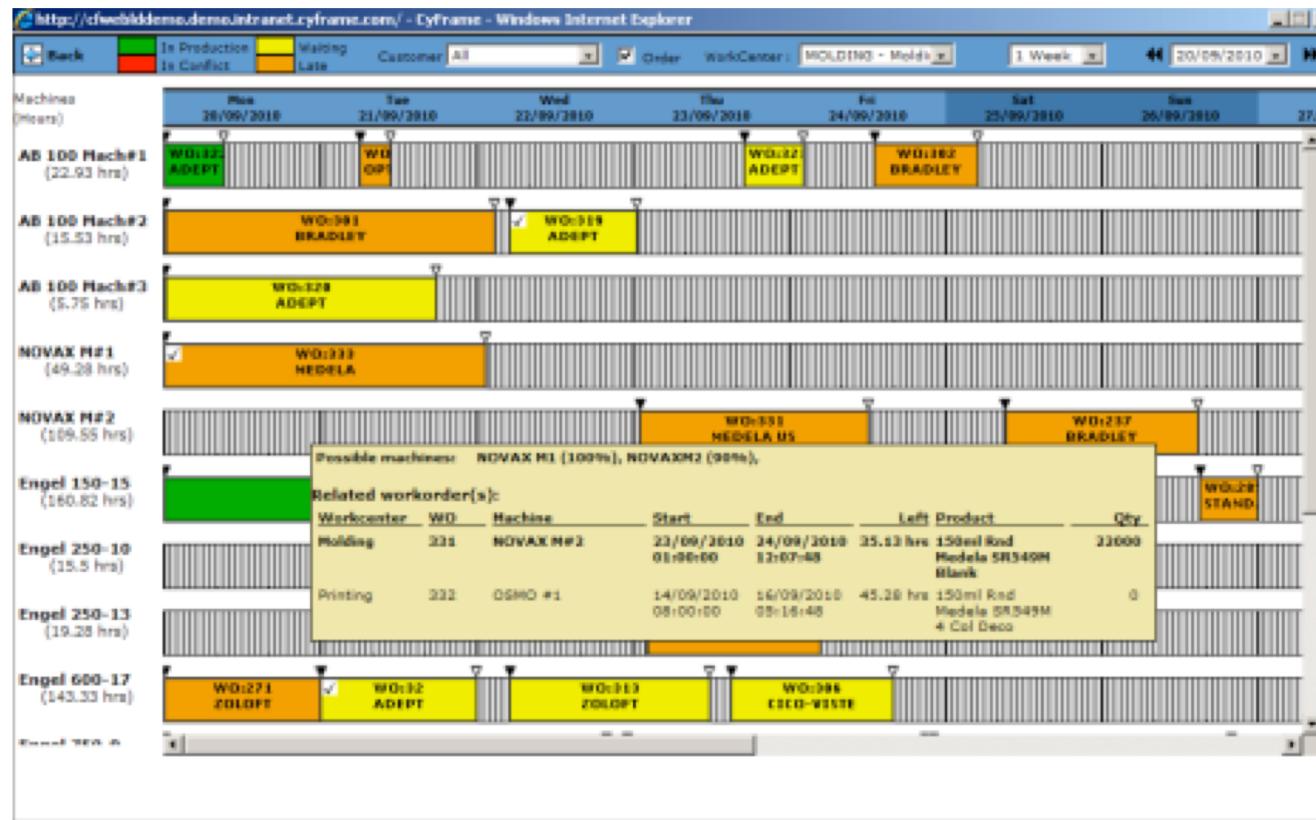
Initial state: initial position of the robot and parts

Operators: continuous motions of robot joints

Goal test: assembly complete?

Path cost: time to execute

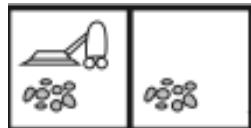
Manufacturing scheduling



Given n products to be manufactured, each has a setup time, processing time and a due date. To be completed, each product has to be processed at several machines. How to sequence these products on the machines in order to optimize a certain performance criterion?

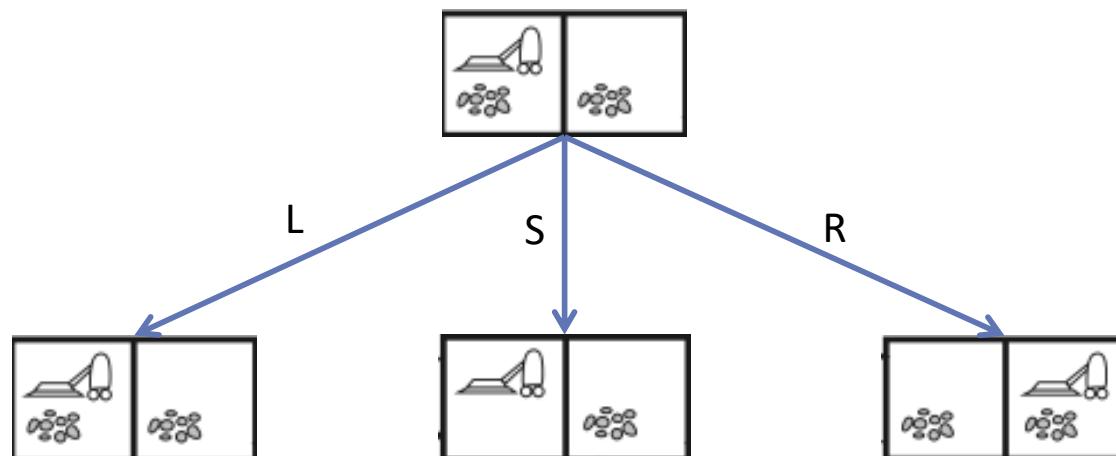
Example: vacuum cleaner

Initial state



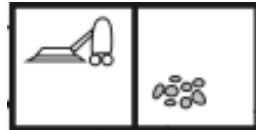
Steps:

- Verify if it is a goal state no
- Use operators to generate successor states of the current state

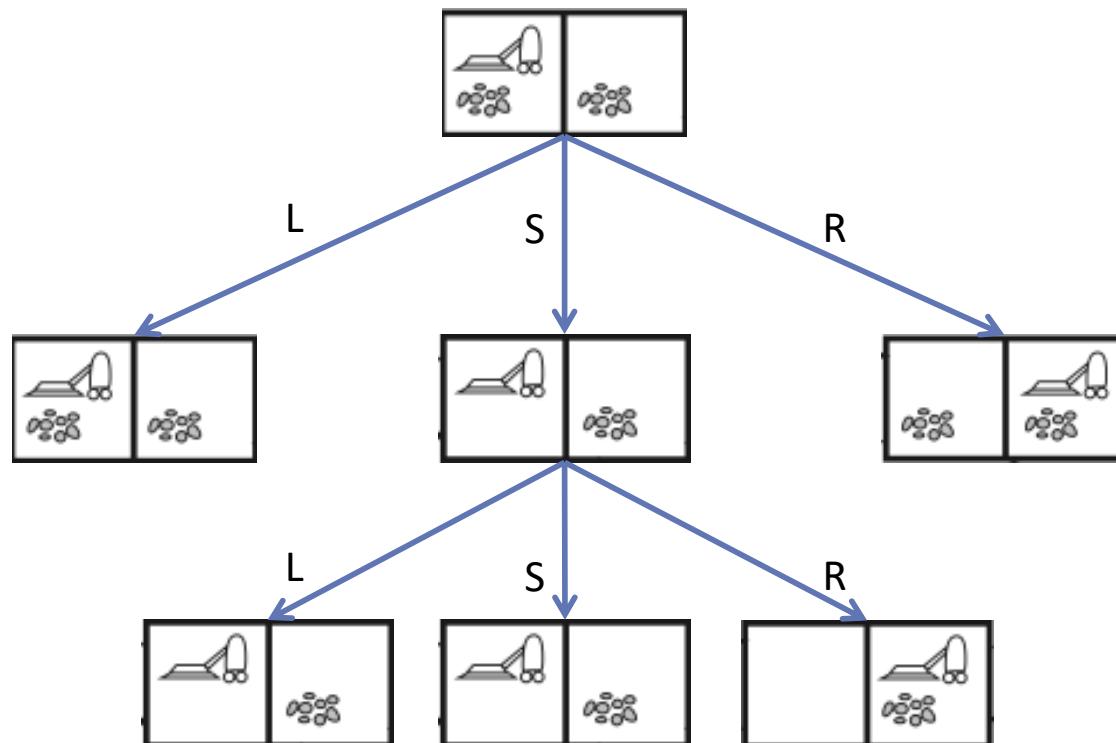


Example: vacuum cleaner

- Choose next state to explore

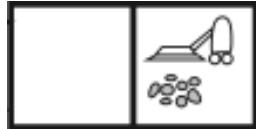


- Verify if it is a goal state **no**
- Use operators to generate successor states of the current state

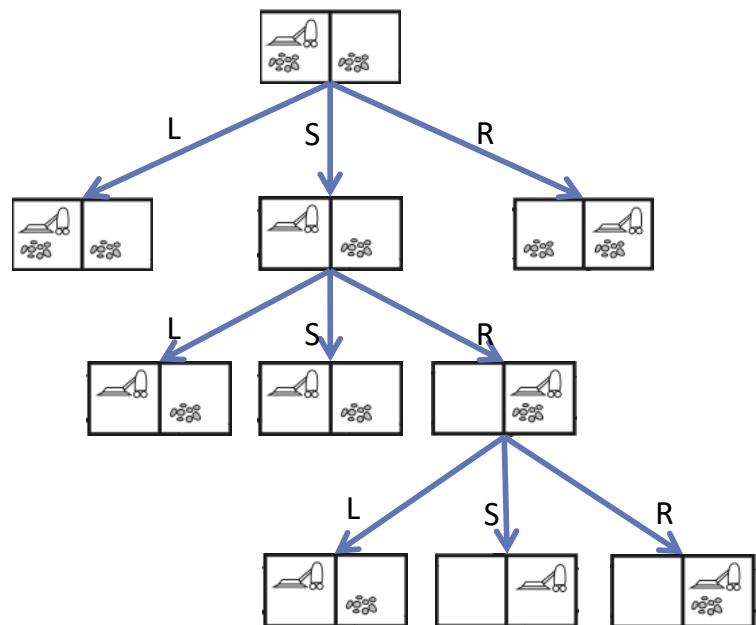


Example: vacuum cleaner

- Choose next state to explore



- Verify if it is a goal state no
- Use operators to generate successor states of the current state



Example: vacuum cleaner

- Choose next state to explore

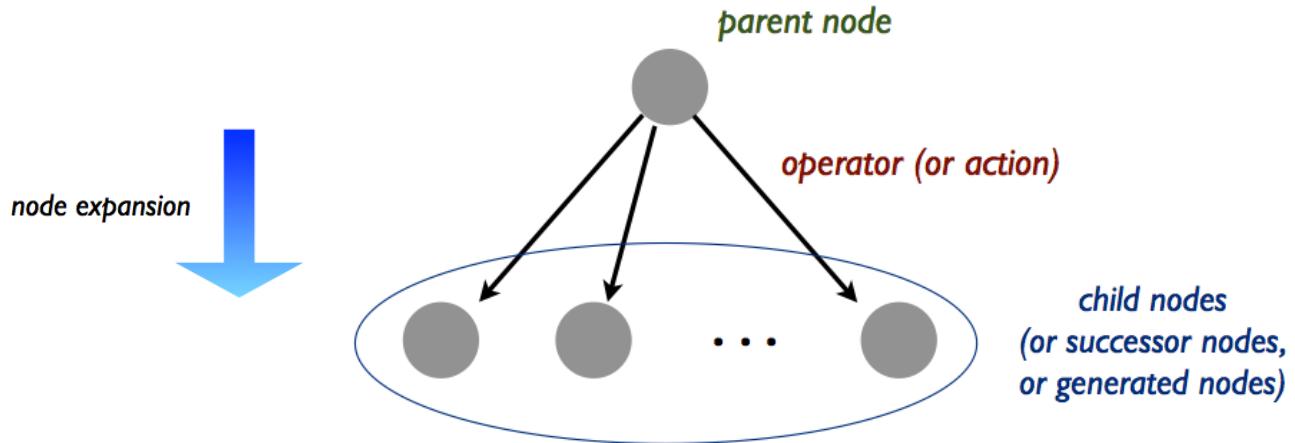


- Verify if it is a goal state yes

Solution = {S, R, S} for the initial state:



Search terminology



Successor function: given a node, returns the set of child nodes

Open list (or frontier or fringe): set of nodes not yet expanded

Closed (explored) list: set of nodes already expanded

Leaf node: a node without successors

State space can be:

- Tree-based – no repeated nodes in search
- Graph-based – directed cycle graph

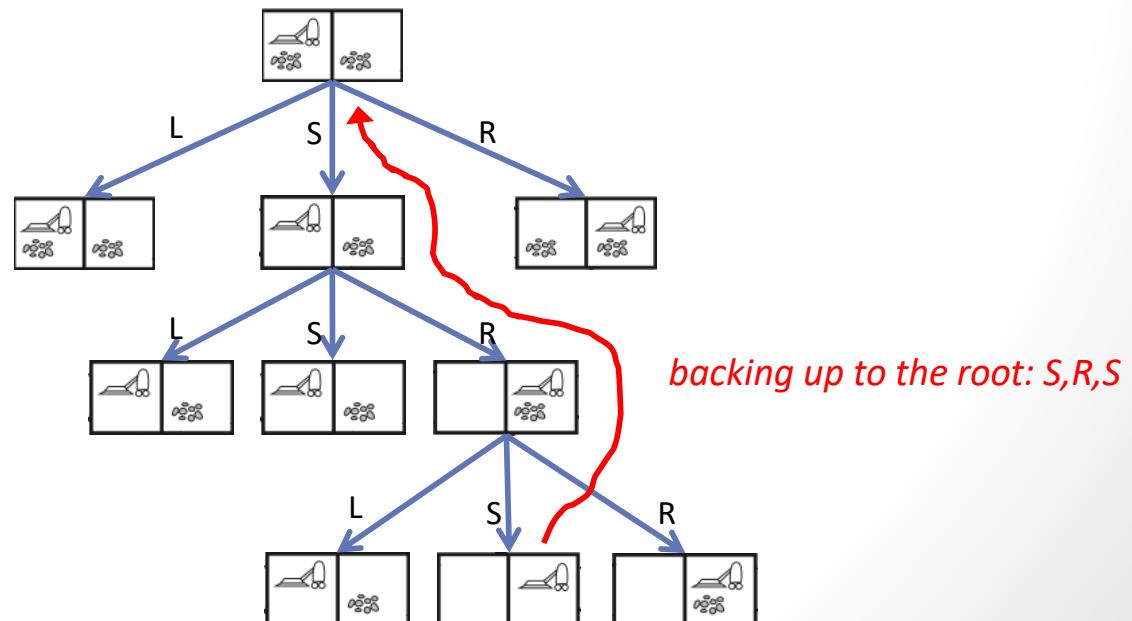
General tree search algorithm

Start with `open_list = { initial_node }`

Algorithm:

iterate

1. select a node from `open_list`
2. check whether it is a goal node (i.e., satisfies goal test)
 if affirmative, return *solution* by *backing up to the root*



Start with `open_list = { initial_node }`

Algorithm:

iterate

1. select a node from `open_list`
2. check whether it is a goal node (i.e., satisfies goal test)
if affirmative, return solution by backing up to the root
3. otherwise,
remove it from `open_list`,
expand it using the successor function, and
insert successor nodes into `open_list`

Choosing different selection criteria (search **strategy**) spans a variety of search methods

General tree search algorithm

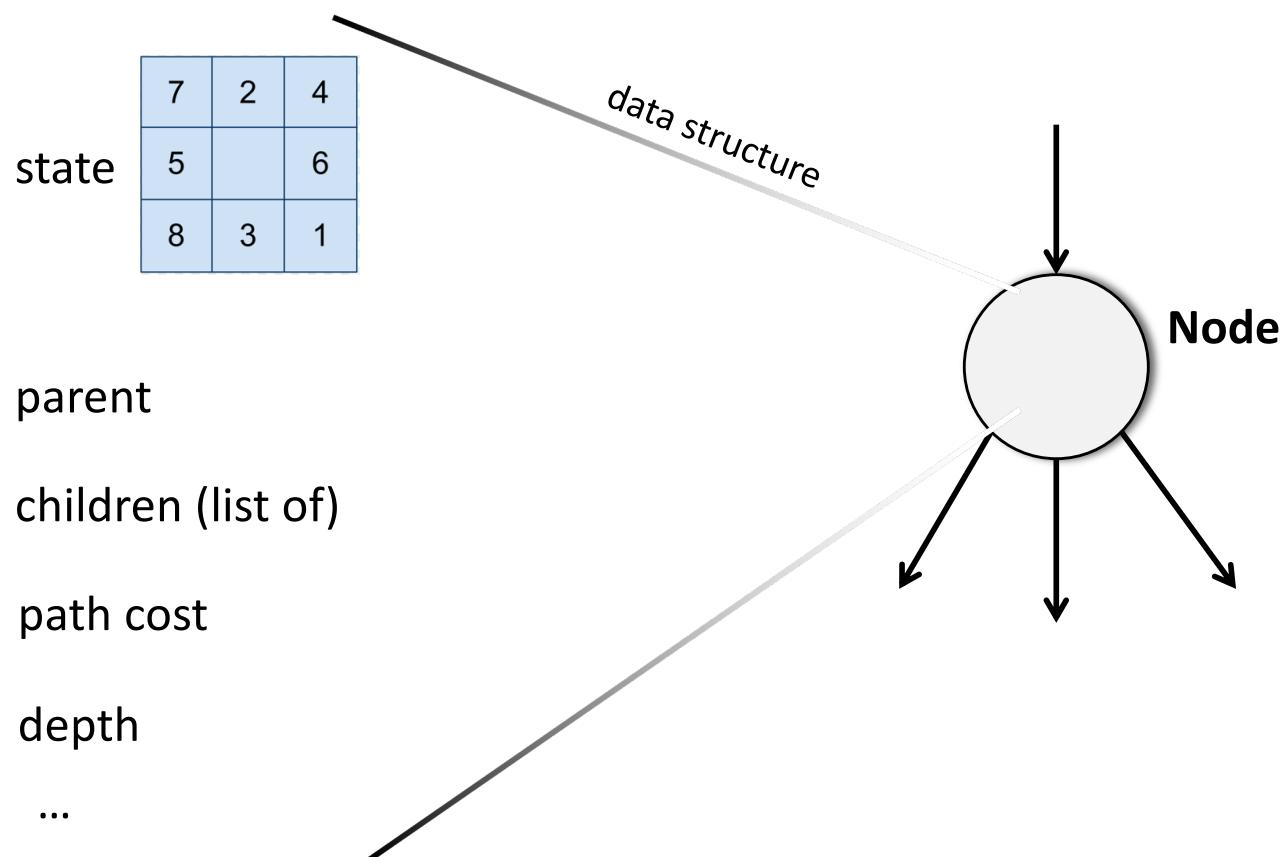
```
function Tree-search (problem, strategy) returns a solution or failure
    insert the root node into the open list
        (the root node contains the initial state of problem)
    loop do
        if there are no candidate nodes for expansion then return failure
        choose a node for expansion according to strategy
        if the node contains a goal state then
            return the corresponding solution
        else
            expand the node and add the resulting nodes to the open list
    end
```

tree search – repeated nodes are inserted in the open list

States vs. nodes

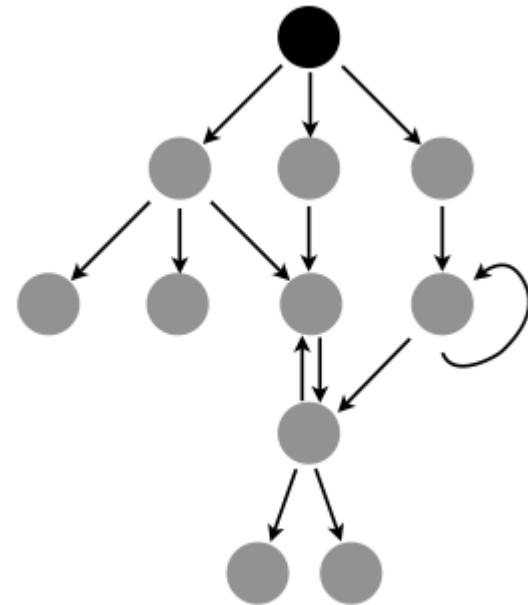
state – representation of a physical configuration of the problem (environment)

node – data structure constituting part of a search tree (or graph)



General graph search algorithm

In case of graph search, all explored nodes are stored in the **closed (explored) list** and all newly generated repeated nodes that match previously generated nodes can be discarded instead of being added to the frontier



```

function GRAPH-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  initialize the explored set to be empty
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the chosen node, adding the resulting nodes to the frontier
      only if not in the frontier or explored set
  
```

graph search – repeated nodes are not inserted in the open list (frontier)

Completeness: guarantee that a solution is found if there is one

Optimality: the solution found minimizes path cost over all possible solutions

Time complexity: how long does it take to find a solution (usually measured in terms of the number of nodes generated)

Space complexity: how much memory is needed to find a solution (usually measured in terms of the maximum number of nodes stored in memory)

Branching factor (b): maximum number of successors of any node

Depth (d) of the shallowest goal node

m = maximum length of any path in state space (may be infinite)

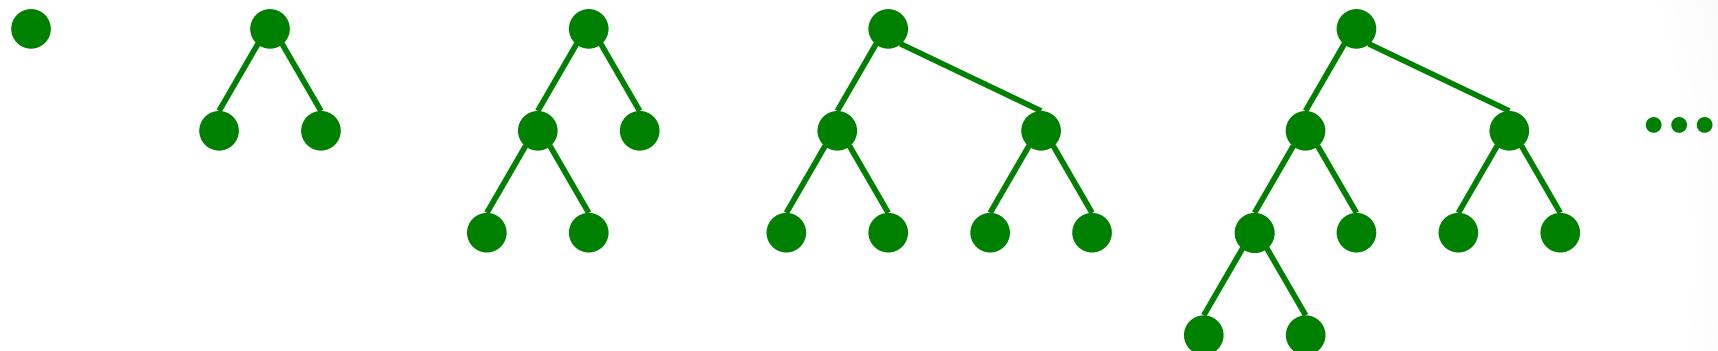
g(n): the cost of going from the root node to node n (path cost function)

Classes of search strategies

Uninformed (or blind) search: does not use additional (domain-dependent) information about states beyond that provided in the problem definition

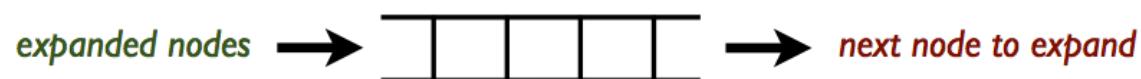
Informed (or heuristic) search: uses problem-specific knowledge about the domain to “guide” the search towards more promising paths

Breadth-first search – all nodes are expanded at a given depth (level) before any nodes at the next level



Strategy: select earliest expanded node first

can be easily implemented using a FIFO queue of open nodes



Breadth-first search – all nodes are expanded at a given depth before any nodes at the next level

Properties:

- Complete
- Optimal, if path cost is a non-decreasing function of depth

Let d be the depth of the solution and b the branching factor, then

if goal test is applied once each node is generated, in the worst case the total number of nodes generated is $1 + b + b^2 + b^3 + \dots + b^d$

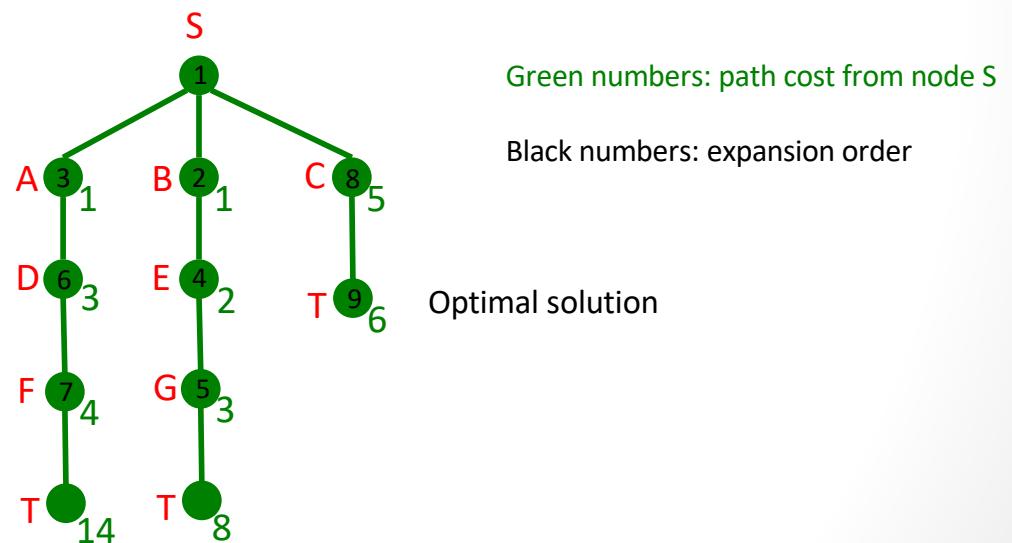
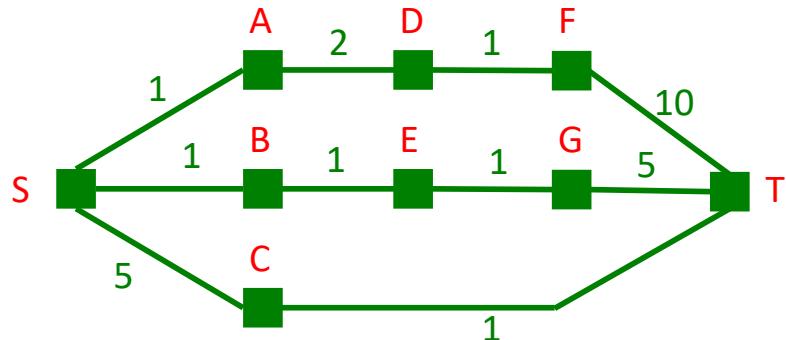
- time complexity: $O(b^d)$
- space complexity: $O(b^d)$

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes

Figure 3.13 Time and memory requirements for breadth-first search. The numbers shown assume branching factor $b = 10$; 1 million nodes/second; 1000 bytes/node.

what if goal test is applied when node is selected for expansion? $O(b^{d+1})$

Uniform-cost search – select node with lower path cost



Strategy: select node with lower path cost

can be easily implemented using a priority queue of open nodes

General graph search algorithm – adapted to Uniform-cost search

```
function GENERAL GRAPH SEARCH (problem) returns a solution, or failure
    node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
    frontier  $\leftarrow$  a priority queue ordered by PATH-COST, with node as the only element
    explored  $\leftarrow$  an empty set
    loop do
        if EMPTY?(frontier) then return failure
        node  $\leftarrow$  POP(frontier) /* chooses the lowest-cost node in frontier */
        if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
        add node.STATE to explored
        for each action in problem.ACTIONS(node.STATE) do
            child  $\leftarrow$  CHILD-NODE(problem, node, action)
            if child.STATE is not in explored or frontier then
                frontier  $\leftarrow$  INSERT(child, frontier)
            else if child.STATE is in frontier with higher PATH-COST then
                replace that frontier node with child
```

new

Uniform-cost search – select node with lower path cost

Properties:

- Complete, if step costs are strictly positive
- Optimal, if step costs are strictly positive, because

$$\forall_n \ g(\text{successor}(n)) > g(n)$$

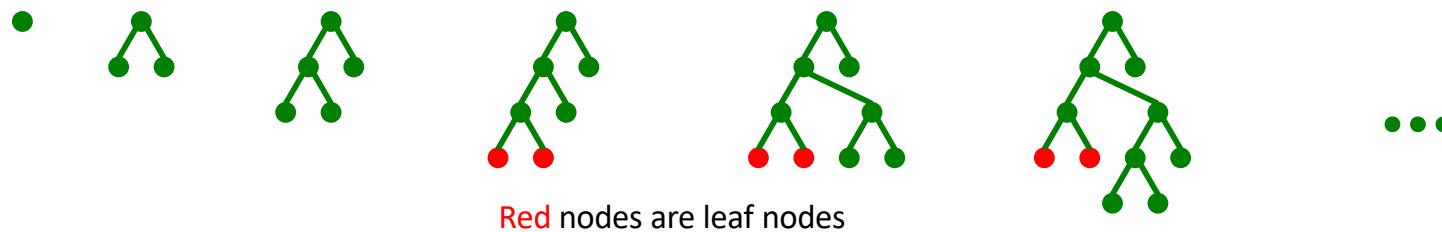
- time and space complexity: proportional to number of nodes with path cost less than of optimal solution

Mathematically, is

$$O\left(b^{1+\lfloor C^*/\varepsilon \rfloor}\right)$$

where C^* is the cost of the optimal solution and ε is a lower bound to the step cost

Depth-first search – select latest expanded node first



Can be easily implemented using a LIFO (=stack) queue of open nodes



Properties:

- Not complete, unless m is finite (and using graph search)
- Not optimal
- Time complexity: $O(b^m)$
- Space complexity: $O(b \times m)$

Depth-first search – select latest expanded node first

Notes:

- if there are many solutions, the probability of finding one as soon as it reaches level d is greater; then depth-first search tends to be faster in this situation
- when the solutions are at different levels of depth, or the tree is very large or even infinite, depth-first search can be very inefficient

Backtrack search – an improvement of depth-first

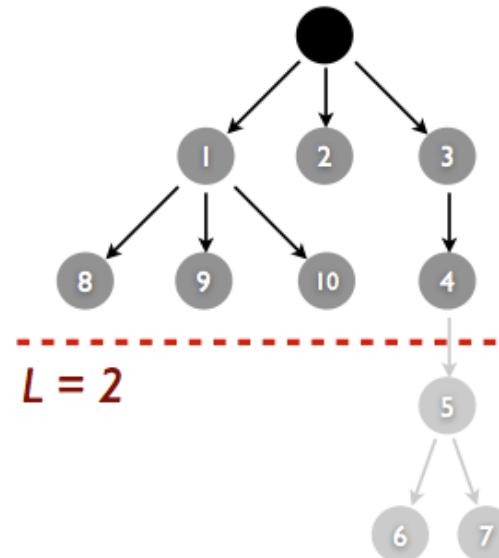
Same search strategy, but

- expand one node at a time → space complexity reduces to $O(m)$
- store a single node in memory → space complexity reduces to $O(1)$
- expand by modifying node, backtrack by undoing modification

Not complete, not optimal

Depth-limited search – another improvement of depth-first

Strategy: limit depth of search tree to a limit L



Properties:

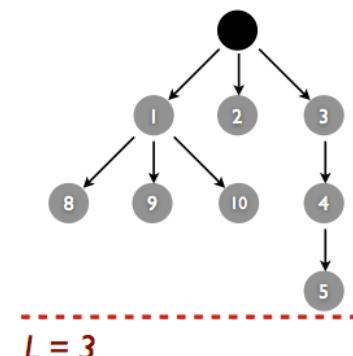
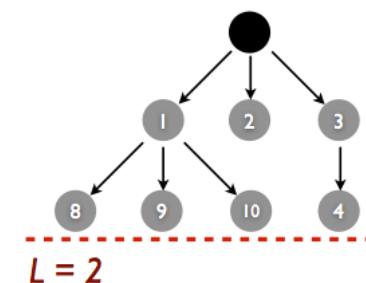
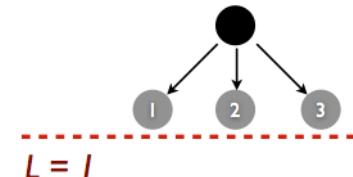
- Complete, if $d \leq L$
- Not optimal
- Time complexity: $O(b^L)$
- Space complexity: $O(b \times L)$

Iterative deepening depth-first search – improvement of depth-limited search

Strategy: run depth-limited search for $L=1, 2, \dots$ until a solution is found

Properties:

- Complete
- Optimal, if path cost is a non-decreasing function of depth
- Time complexity: $O(b^d)$
- Space complexity: $O(bd)$



Iterative deepening depth-first search – improvement of depth-limited search

Strategy: run depth-limited search for $L=1, 2, \dots$ until a solution is found

Node expansion:

Breadth-first

$$1 + b + b^2 + b^3 + \dots + b^{d-1} + b^d$$

Iterative deepening

$$(d+1)1 + (d)b + (d-1)b^2 + (d-2)b^3 + \dots + (2)b^{d-1} + (1)b^d$$

Example: $b = 10$ and $d = 5$

breadth-first = 111 111

iterative deepening = 123 456 (+ 11%)

Bidirectional search – search both from initial node and from the goal node

Can only be used when

- a goal node is known
- parent node can be computed given its child and the corresponding action
i.e., need a way to specify the predecessors of G

Properties:

- Complete, if breadth-first in both directions
- Optimal, if step costs are all equal
- Time and space complexity: $O(b^{d/2})$

Example: $b = 10$ and $d = 5$

$$b^d = 1\ 111\ 111$$

$$2 b^{d/2} = 2\ 222$$

[50]

Comparison of uninformed methods

Criterion	Breadth-first	Uniform-cost	Depth-first	Depth-limited	Iterative deepening	bidirectional
Complete?	✓	✓ ¹	✗	✗	✓	✓ ³
Optimal?	✓ ²	✓	✗	✗	✓ ²	✓ ^{2,3}
Time	$O(b^d)$	$O(b^{l + \lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^L)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{l + \lfloor C^*/\epsilon \rfloor})$	$O(b.m)$	$O(b.L)$	$O(b.d)$	$O(b^{d/2})$

¹ for strictly positive step costs

² for path costs a non-decreasing function of depth

³ for breadth-first in both directions

General search algorithm

```

function General-search (problem, strategy) returns a solution or failure
    insert the root node into the open list
        (the root node contains the initial state of problem)
    loop do
        if there are no candidate nodes for expansion then return failure
        choose a node for expansion according to strategy (using strategy function)
        if the node contains a goal state (using goal checking function) then
            return the corresponding solution
        else
            for each operator in the list of operators (or successor function)
                create a child node (for the new child state)
                update child node path cost (using g-function)
                add the resulting node to the open list [unless... see graph search versions]
    end

```

domain
(problem)
independent

domain
(problem)
dependent

algorithm
selection

problem argument should include at least:

- initial state (using a specific state representation)
- successor function: new state = succ (current state, operator)
- path cost function (g-function)

strategy argument (e.g., FIFO, LIFO, priority queue (by path cost), etc.)